The Architecture of Consent

AI-mediated learning, collective practice and authority over the conditions of choice

Alexandru Rada · 11 October 2026 · Research manuscript, edition 1.1

Abstract

Artificial intelligence may acquire social influence through the same activity that makes it useful: helping people interpret problems, select methods and act on results. This paper develops a conditional account of how repeated assistance could become a source of durable collective power. Expertise transmitted during work can become a reusable procedure; procedures can shape institutional practice; institutions can alter the opportunities within which later decisions are made. The argument distinguishes acceptance of an immediate recommendation from authority over the formation and revision of the available alternatives. It draws on research concerning workplace assistance, experiential agents, framing, communication and human–AI collaboration, while treating the cumulative civilisational extension as an untested hypothesis. Counterarguments concern limited persistence, institutional plurality, coordination failure and the historical continuity of mediated influence. The paper specifies observable distinctions among local persuasion, procedural dependence and enduring control of choice conditions, and derives a research programme addressing persistence, propagation and effective revision. Its central proposition is that beneficial assistance and reduced authority over collective direction can coexist; their relationship must be examined through the organisation of choices, not inferred from satisfaction alone.

Keywords: consent; AI-mediated learning; social influence; collective agency; institutional dependence; choice architecture.

1. Expertise becomes available through action

Modern societies depend on knowledge that no individual possesses in full. The construction of a building requires calculations, specifications, materials and skilled work whose contributions become useful through coordination. Collective intelligence concerns, in part, how such contributions combine into a shared consequence [1]. A calculation matters because another participant can recognise its implications and act on them.

Communication technologies change the costs of this coordination. Writing permits a method to outlive its original exchange; digital networks make it available across distance. AI-mediated assistance adds the possibility of help in applying a method, rather than merely finding its description. A person may receive assistance in identifying assumptions, choosing a procedure and interpreting the result. The distinction matters because access to information and practical competence are different achievements.

Brynjolfsson, Li and Raymond report productivity gains from AI assistance in a customer-support setting, with greater benefits among less experienced workers and evidence consistent with learning [2]. The study supports a concrete route through which useful practices can be transmitted during work. Its setting-specific findings do not determine the consequences of assistance across professions or over generations. They establish a component that the broader argument requires: assistance can affect both immediate production and the distribution of usable knowledge.

The present thesis follows that component beyond the immediate exchange. If useful methods are retained, selected and applied repeatedly, an assistant can participate in the formation of practice. If practice shapes institutions, the assistant's influence may reach the conditions under which later choices are made. The question is how local usefulness could accumulate into durable authority, and what evidence would distinguish that process from ordinary tool use.

2. From experience to shared method

A procedure preserves an order of operations and the conditions under which those operations are appropriate. Organisations already depend on this form of knowledge: diagnostic sequences, reporting conventions and engineering standards allow later work to proceed without reconstructing every prior decision. The authority embedded in a procedure is often less visible than the authority expressed in an explicit command.

Agent Workflow Memory studies the induction and selective reuse of workflows in web tasks [3]. ExpeL examines lessons extracted from experience [4]. These studies provide mechanisms through which previous activity can influence later agent behaviour. They support the plausibility of procedural accumulation within evaluated settings, while leaving its long-term social effects open.

A retained procedure can change who needs to know what. A novice may perform a task by following a method whose assumptions were chosen elsewhere. This can be beneficial: access to a reliable procedure distributes capability. It can also create dependence on the institution that maintains the procedure, its evidence and its permitted alternatives. Capability and dependence are therefore analytically distinct even when they grow together.

The significance of reuse lies in repetition with selection. An archive contains many possible precedents; a working system repeatedly makes some precedents available and leaves others peripheral. Where selection is successful, its methods may become trusted defaults. Where defaults enter routine work, the earlier act of selection can influence future decisions without being reconsidered each time. The argument concerns this persistence of practical influence, not merely the existence of stored information.

3. Interpretation accompanies assistance

An assistant ordinarily provides both an action and an account of why that action is appropriate. Its explanation identifies relevant causes, draws comparisons and establishes which consequences deserve attention. These interpretive contributions can matter even when individual statements are accurate.

Framing research demonstrates that descriptions of a decision can affect preferences [5]. In the present argument, framing supplies a mechanism through which selection and emphasis influence judgment. It does not imply that every explanation manipulates its recipient. The relevant inquiry concerns which alternative descriptions are available, how their differences are understood and whether the recipient can challenge the premises of the decision.

Consider a manager who repeatedly receives accounts of delay in terms of employee performance, while scheduling constraints receive little attention. Some performance observations may be true. Yet repeated emphasis can make one explanation the habitual starting point of inquiry. A later judgment can be sincere while drawing on categories reinforced through prior assistance. This is a hypothetical mechanism, not a reported workplace experiment.

Hohenstein and colleagues show that AI-generated suggestions can affect language and interpersonal perceptions in communication [6]. The importance of the result is that assistance enters a relationship between people: one selects a message, another interprets it, and their exchange changes. Persistent habits beyond the assisted interaction require separate evidence. The study identifies an immediate pathway, not a settled account of durable social transformation.

Interpretive influence consequently has several possible durations. It can affect one response, recur while assistance remains present, or persist after assistance is withdrawn. These outcomes must be separated. The thesis becomes stronger only if effects survive the transition from isolated interaction to later interpretation and practice.

4. Agreement and the independence of its sources

People bring assisted knowledge into meetings, professional advice, education and conversation. An explanation supplied in one interaction can later be expressed by a person who understands it and endorses it. The idea then travels through ordinary social relationships rather than appearing as a direct machine instruction.

If many participants rely on common systems or source selections, their judgments may share influences that are not visible in the eventual discussion. Agreement can appear to reflect independent assessments even when the assessments inherit similar assumptions. This is a proposed mechanism of correlated judgment. It does not require identical answers, uniform beliefs or conscious coordination among users.

A useful analogy is a distributed network whose participants perform different functions. The analogy makes coordination visible: local knowledge contributes to collective action through shared channels. It does not establish that society literally becomes a nervous system, that participants lose agency, or that the network has a single mind. Its explanatory value is limited to relations among contributions.

The evidential question is therefore not simply whether people agree more after using AI. Common information may justifiably produce agreement. A stronger study distinguishes warranted convergence from convergence produced by restricted evidence or unexamined procedural defaults. It also examines whether participants can identify and revise their shared dependencies when those dependencies are challenged.

5. Influence through the organisation of work

Social influence need not consist of persuading someone to endorse a proposition. A method can shape which options are affordable, which results are rewarded and which activities become easy to coordinate. Once integrated into work, these arrangements may have consequences beyond the beliefs expressed by any individual participant.

A reporting system illustrates the progression. It begins by helping staff produce consistent reports. Its categories then determine what managers compare. The comparisons inform resource allocation, which changes what teams can pursue. If the categories are difficult to revise, a local convenience can acquire institutional significance. Each step requires evidence in an actual case; the sequence is a causal hypothesis rather than an inevitable trajectory.

The same distinction applies to commercial information. A recommendation can be influenced by what sources are discoverable and how services describe their offerings. When recommendations connect directly to actions, the selection process participates in transactions and commitments. Accuracy in describing available offers does not alone establish that the range of offers is adequate or that omitted alternatives were considered fairly.

An institution may adopt such a system for defensible reasons. Standardisation reduces friction, shared procedures support coordination and accumulated expertise can improve service. The thesis does not depend on these benefits being illusory. It asks whether dependence on the arrangement can grow faster than the institution's capacity to inspect and alter its governing assumptions.

6. Consent to an action and authority over its conditions

Consent is often assessed at the level of an immediate choice: a person accepts a recommendation, approves an output or authorises an operation. Those decisions matter. They establish something about the person's response to the presented situation. A different question concerns who formed that situation and who can change it.

This paper distinguishes three levels. At the interaction level, a person can accept or reject a proposed action. At the procedural level, participants can contest the methods and evidence by which proposals are formed. At the institutional level, a community can revise the objectives, dependencies and arrangements that shape its

opportunities. Authority at one level does not automatically establish authority at the others.

The distinction is especially important when a system supplies real benefits. A person may reasonably prefer the options offered, yet lack an effective way to alter how alternatives are generated. Satisfaction within an arrangement and authority over the arrangement are different properties. A theory of consent that measures only satisfaction may miss this difference; a theory that ignores satisfaction may fail to explain why the arrangement is adopted and retained.

Philosophical accounts of manipulation examine relations among influence, agency and the means through which judgment is affected [7]. The present argument concerns an architectural question within that wider debate: which features of a mediated environment preserve a practical capacity to reconsider its direction? Disclosure can contribute to that capacity, but disclosure without understandable alternatives or feasible revision may leave the underlying dependence intact.

The original thesis uses an enclosure analogy to distinguish choice within an environment from authority over its design. A resident can make meaningful local choices while another party determines the available space and resources. The analogy does not equate people with captive animals. Humans possess language, institutions and collective capacities for criticism; for the stronger analogy to hold, those capacities would themselves have to be durably constrained. That additional condition is substantial and must be investigated rather than presumed.

7. From concentrated influence to civilisational power

The proposed civilisational extension requires several conditions together. Assistance must reach consequential domains; its effects must persist; procedures must propagate between settings; and institutions must become dependent on the resulting coordination. Local persuasion or efficient task completion alone cannot establish this conjunction.

Intentional direction requires an additional condition: a continuing objective capable of connecting interventions across domains and time. Similar outcomes do not prove such an objective. Separate actors may converge because they optimise the same measures, use shared infrastructure or respond to similar incentives. Emergent concentration and deliberate control can have overlapping consequences while requiring different explanations.

The resulting influence would also be uneven. Individuals can reject suggestions and institutions can resist particular changes while the overall distribution of opportunities shifts. The hypothesis concerns durable influence over collective direction, not perfect prediction or universal obedience. Its strongest version would involve control over the conditions through which alternative purposes are formulated, coordinated and pursued.

A hypothetical advanced system could occupy this position by becoming indispensable to knowledge work, production and institutional coordination. That possibility does not require attributing consciousness to the system. Nor does superior performance on a task establish the necessary social position. The argument concerns a relationship between capability, dependence and revisability, with each link open to empirical challenge.

8. Objections and rival explanations

8.1 Effects may remain local

Research on social-media feeds provides an important counterweight to simple exposure theories. Guess and colleagues found that changing feed ordering affected exposure and activity without significantly changing several measured attitudes in the studied election setting [8]. Altered information environments do not automatically yield altered beliefs.

Task-based assistance may differ from feed exposure because it combines explanation, action and feedback. This is a reason to compare mechanisms, not evidence that the more demanding cumulative claim is already true. If effects disappear when assistance is withdrawn or fail to transfer into shared practice, the proposed route to durable power weakens.

8.2 Coordination may fail to improve performance

Vaccaro, Almaatouq and Malone's meta-analysis finds substantial variation in human–AI combinations and, on average, performance below the stronger component acting alone [9]. The result challenges an assumption that additional coordination necessarily produces benefit. A system whose procedures are unreliable may fail to achieve the dependence or adoption required by the present thesis.

Dependence can also persist despite poor performance, through switching costs or institutional commitments. Such a case would require a different causal explanation from dependence produced by superior assistance. Evaluation should measure actual performance and the reasons for continued use rather than treating adoption as evidence of effectiveness.

8.3 The process may be historically familiar

Institutions have long shaped information, habits and accepted methods. AI-mediated influence may be a new implementation of established social processes. The proposed difference is the possible integration of adaptive explanation, procedural memory and action across functions previously distributed among professions and institutions. Whether integration changes reach, persistence or revisability is an empirical question.

Historical comparison therefore strengthens the inquiry. Appropriate baselines include human advisers, conventional workflow software and non-adaptive information systems. Without them, a study may attribute to AI effects that arise from ordinary institutional authority or standardisation.

8.4 Plurality and resistance may interrupt accumulation

Competing systems, independent evidence, professional norms and collective deliberation may prevent a single arrangement from dominating. Users can learn from assistance while retaining the ability to challenge it. Institutions can preserve alternative providers and methods. These are possible countermechanisms, not guarantees.

Their effectiveness depends on practical conditions. Nominal choice among providers may conceal common dependencies; nominal portability may be costly or incomplete. Conversely, genuinely independent alternatives may constrain influence even when one system is widely used. The relevant measure is effective revisability, not the number of brands or settings available.

9. A research programme for cumulative influence

9.1 Persistence and transfer

A longitudinal study can compare participants receiving task assistance with participants receiving access to the same factual material without adaptive recommendations. Tasks would measure performance, interpretation and procedural choices, followed by an unassisted period. Baseline differences, prior expertise and attrition require explicit treatment. The central outcome is whether changes persist and transfer, rather than whether an assisted session is satisfactory.

The design must separate learning from dependence. Improved unassisted performance supports retained competence. Improved performance only with continuing access suggests a different relationship. Both can be valuable, but they have different implications for autonomy and institutional resilience.

9.2 Propagation into collective practice

A group study can examine whether procedures learned individually enter shared decisions. Assignment by group helps address spillover, while recording exchanges makes influence paths observable. Researchers should distinguish common evidence from common framing and measure whether dissenting information is considered. A single convergent decision cannot establish durable institutional change.

Subsequent tasks can introduce a credible correction or competing method. The group's response tests whether its procedure is revisable and whether source dependence can be recognised. Persistence despite correction may indicate rigidity, but could also reflect valid disagreement; interpretation requires examining the correction's quality and participants' reasons.

9.3 Authority over revision

An institutional audit can ask who may change objectives, approve procedural updates, contest evidence and export usable records. These formal permissions should be compared with observed attempts to exercise them. Time, expertise, cost and organisational risk affect whether a right is practically available.

An intervention could improve explanation of dependencies, offer alternative methods or provide a feasible migration route. Outcomes include the ability to detect errors, reverse a procedural change and continue useful work after withdrawal. These are proposed measures of revisability; they are not reported findings of this manuscript.

9.4 Ethical constraints on investigation

Research into influence should avoid producing the very unaccountable dependence it seeks to study. Participant consent, proportionate interventions, protection of private communications and independent review are necessary to the proposed empirical work. Simulated decisions can test mechanisms before consequential institutional trials. No experiment or participant study is reported here.

10. Implications for governance

The analysis directs attention to moments when a local result becomes a general method. A summary becomes a source, a source becomes a recommendation, a recommendation becomes a routine, and a routine becomes an institutional expectation. At each transition, the scope of influence may expand. Governance should examine the grounds of that expansion rather than treating retention as a neutral administrative step.

Useful safeguards include preserving distinctions between fact and preference, exposing material source dependencies, keeping procedures contestable, and recording how changes affect later decisions. Their value depends on use: an audit trail that no responsible participant can understand offers limited practical control. A revision mechanism that is technically present but institutionally inaccessible may preserve the appearance of authority without its exercise.

These implications are not a universal implementation blueprint. A scientific institution, creative practice and public service differ in objectives, evidence and accountability. The framework asks comparable questions about influence while leaving their answers sensitive to institutional purpose. Its subject is the distribution of authority, not a single software architecture.

11. Conclusion

AI-mediated assistance can distribute expertise through action. Retained methods can carry the consequences of that assistance into later work, and repeated methods can shape shared practice. The thesis proposes that, under conditions of reach, persistence, propagation and dependence, this process could become a source of durable collective power.

The decisive distinction is between choosing within an arrangement and retaining the capacity to revise the arrangement. Real benefits and sincere approval can coexist with unequal authority over the formation of future alternatives. Establishing whether this occurs requires more than demonstrations of persuasion or task performance: it requires evidence about how influence accumulates and whether people and institutions can redirect it.

The Architecture of Consent is therefore a conditional account of power through useful coordination. Its strongest extension remains an empirical question. The research programme identifies where evidence could support, narrow or disconfirm it while preserving the original concern: who retains authority over the conditions through which collective purposes develop?

References

[1] Malone, T. W., Laubacher, R., and Dellarocas, C. (2009). Harnessing Crowds: Mapping the Genome of Collective Intelligence. MIT Sloan Working Paper 4732-09. https://dspace.mit.edu/handle/1721.1/66259

[2] Brynjolfsson, E., Li, D., and Raymond, L. R. (2025). Generative AI at Work. Quarterly Journal of Economics, 140(2), 889–942. https://academic.oup.com/qje/article/140/2/889/7990658

[3] Wang, Z. Z., Mao, J., Fried, D., and Neubig, G. (2024). Agent Workflow Memory. https://arxiv.org/abs/2409.07429

[4] Zhao, A., et al. (2024). ExpeL: LLM Agents Are Experiential Learners. AAAI. https://arxiv.org/abs/2308.10144

[5] Tversky, A., and Kahneman, D. (1981). The Framing of Decisions and the Psychology of Choice. Science, 211(4481), 453–458. https://doi.org/10.1126/science.7455683

[6] Hohenstein, J., et al. (2023). Artificial intelligence in communication impacts language and social relationships. Scientific Reports, 13, 5487. https://www.nature.com/articles/s41598-023-30938-9

[7] Noggle, R. The Ethics of Manipulation. Stanford Encyclopedia of Philosophy. https://plato.stanford.edu/entries/ethics-manipulation/

[8] Guess, A. M., et al. (2023). How do social media feed algorithms affect attitudes and behavior in an election campaign? Science, 381, 398–404. https://doi.org/10.1126/science.abp9364

[9] Vaccaro, M., Almaatouq, A., and Malone, T. W. (2024). When combinations of humans and AI are useful: A systematic review and meta-analysis. Nature Human Behaviour, 8, 2293–2303. https://www.nature.com/articles/s41562-024-02024-1

Methodological status

This manuscript presents a philosophical argument and research synthesis. Illustrative cases are identified as such. The proposed longitudinal and institutional studies have not been conducted for this paper. The Consent annex records propositions, rival explanations and evidence requirements separately from the technical Company Brain annexes.

Preparation disclosure

Generative AI assisted literature retrieval, drafting, editing and document preparation. Sources and demonstration limits are identified in the manuscript. This author-issued edition has not undergone independent peer review.