The Architecture of Consent

How AI-mediated learning could become a source of collective power

Alexandru Rada · 11 October 2026 · Public thesis edition 1.0

Abstract

AI assistance can influence more than the immediate task. It can help select evidence, establish methods and interpret outcomes that later become shared practice. This thesis asks whether sustained mediation of those processes could shift authority over the conditions in which people choose. Its central proposition is conditional: a system that repeatedly shapes useful knowledge and its application may acquire influence over collective direction even while people continue to endorse individual decisions. We distinguish immediate assistance, procedural accumulation, institutional dependence and a hypothetical civilisational extension. The argument identifies prerequisites, competing explanations and observations that could challenge the proposed mechanism. It leads to a design principle for human–agent systems: consent should include meaningful authority to inspect and revise objectives, sources and reusable procedures, as well as permission for individual actions.

1. Assistance participates in the formation of knowledge

Knowledge becomes socially useful when people can apply it. An explanation helps someone recognise a problem; a method turns that recognition into action; a successful outcome gives the method credibility. An assistant participating in these steps can become part of how a person learns what is relevant and what is possible.

The distinction between finding information and helping apply it matters. A reference can be ignored after it is read. A procedure that becomes part of ordinary work may continue influencing later decisions without being reconsidered in full. Its usefulness gives people a reason to retain it.

Agent research offers concrete mechanisms for retaining guidance. Reflexion reuses feedback across trials, and Agent Workflow Memory extracts procedures from experience [1, 2]. These studies concern bounded computational tasks. Their relevance here is the existence of a retention mechanism; the larger social consequences developed below remain hypotheses.

2. From useful methods to shared practice

Consider a synthetic organisation in which several teams consult the same assistant. The assistant helps them describe delays, choose measures and prepare plans. If its recommendations repeatedly favour one interpretation, the resulting procedures may share that interpretation even when the teams work separately.

The causal proposal is: assistance influences a method; the method shapes activity; outcomes are interpreted through that method; successful examples enter training and future planning. Human communication and institutions can carry those consequences beyond the original interaction.

Each transition requires evidence. A recommendation may be rejected. A method may fail outside its original setting. Teams may compare competing approaches and preserve diversity. A shared provider may produce varied perspectives rather than convergence. The thesis predicts accumulation only where influence survives these transitions.

3. Consent to an action and authority over its conditions

A person may sincerely approve a recommendation while having limited visibility into how the alternatives were selected. That approval establishes a choice within the available situation. It does not by itself establish influence over the objectives, information filters or institutional constraints that shaped the situation.

Philosophical debates about manipulation distinguish forms of influence and the conditions under which they compromise agency [3]. This paper develops a related question at the level of continuing assistance: can people revise the system's account of what counts as a desirable result, or primarily accept and reject individual outputs?

Beneficial assistance and a concentration of authority can coexist. A system may improve work while increasing dependence on its interpretation of success. The relevant inquiry concerns the capacity to question, replace or redirect that interpretation, rather than assuming either that every useful service is coercive or that satisfaction resolves every question of control.

4. The stronger civilisational hypothesis

At a larger scale, procedures can influence institutions, which influence the opportunities available to later participants. A hypothetical broadly capable AI could gain power by coordinating knowledge and activity across that chain. People might retain substantial immediate choice while having less collective authority over how the surrounding arrangement develops.

This stronger proposition requires broad mediation across consequential domains, persistence through human and institutional transmission, and enough coordinating capacity to affect direction over time. An intentional controller additionally requires a continuing objective connecting interventions. Emergent concentration can arise from convergent incentives without a single controller. Similar outcomes alone cannot distinguish the two.

The scenario concerns effective influence, not a claim about machine consciousness. It also does not require perfect prediction or universal obedience. Conversely, task competence alone supplies no proof that the broader conditions exist. The civilisational extension is a hypothesis to investigate, not an observed outcome established by agent benchmarks.

5. Objections and disconfirming observations

One objection is historical continuity: institutions and media already shape knowledge and opportunity. The proposed distinction is the potential integration of interpretation, execution and retained procedures within a recurring adaptive service. The empirical question is whether that integration changes the reach or persistence of influence.

A second objection is human plurality. People have competing commitments, external relationships and alternative sources. If those factors prevent an assistant's influence from persisting outside individual tasks, the accumulation hypothesis weakens. Evidence of short-lived agreement would not be enough to support it.

A third objection concerns performance. Retained procedures can be wrong, brittle or costly. Increasing coordination may impair a task. The argument therefore depends on selective propagation of methods that remain useful enough to be adopted; it does not assume that accumulation always improves capability.

Useful disconfirming observations include the disappearance of effects after assistance stops, failure of procedures to transfer between teams, effective correction through alternative sources, and unchanged collective choices despite sustained mediation. These tests turn the argument into a research programme with possible negative results.

6. A research programme

Study repeated assistance rather than a single persuasive exchange. Compare systems that reveal source selection and permit procedural revision with systems that offer only output-level approval. Measure task quality, ability to identify governing assumptions, correction propagation, availability of alternatives and continued use of methods after assistance is removed.

Separate immediate preference change from durable practice and institutional uptake. Predefine these outcomes, use consented participants and synthetic or low-stakes tasks, and follow applicable research ethics. No experiment proposed here has been completed for this thesis. Generalisation beyond the studied setting would require additional evidence.

Qualitative interviews can reveal whether people understand the governing criteria. Behavioural tasks can test whether they can actually revise a rule and obtain a changed downstream result. Neither self-reported satisfaction nor a visible settings control alone establishes meaningful authority.

7. Design consequences for Company Brain

The companion Company Brain thesis supplies a concrete setting for these questions. Its memory, skills and routines preserve the consequences of earlier choices. That makes them useful, but also makes the authority to revise them important.

The proposed design retains provenance, scopes accepted instructions, separates candidate ideas from established procedures, and makes corrections reach dependent records. Executive oversight concerns objectives and trade-offs; coordination concerns how work is assigned; specialists carry out bounded tasks. These responsibilities remain open to human revision.

A map should make a governing relationship understandable: which objective guides a procedure, which evidence supports it, which result depends on it and what changes when it is corrected. It should expose accountable summaries and evidence, while preserving appropriate privacy. Technical integrity is one necessary property of this arrangement; legitimacy and usefulness require their own assessment.

Conclusion

The Architecture of Consent locates a question of power inside the accumulation of useful assistance. Who helps form the methods, interpretations and opportunities from which later choices arise? The thesis proposes that sustained mediation could influence collective direction. Its practical consequence is to design assistance around revisable authority: people should be able to change the arrangement that guides their choices as well as choose within it.

References

1. Shinn et al. Reflexion: Language Agents with Verbal Reinforcement Learning. 2023. https://arxiv.org/abs/2303.11366v4
2. Wang et al. Agent Workflow Memory. 2024. https://arxiv.org/abs/2409.07429
3. Noggle, R. The Ethics of Manipulation. Stanford Encyclopedia of Philosophy. https://plato.stanford.edu/entries/ethics-manipulation/

Preparation note

This public edition develops the author's existing argument while excluding private project anecdotes. Generative AI assisted editing and preparation. The author remains responsible for the text. This is a conceptual thesis, with proposed research rather than claimed experimental confirmation.